

PLACEMENT PREDICTION

By: Sakshi Neeraj

This is a study of publicly accessible data regarding on-campus placement of applicants depending on characteristics such as high school graduation percentage and domain of specialisation.

Dataset for practical- <https://www.kaggle.com/benroshan/factors-affecting-campus-placement>
(<https://www.kaggle.com/benroshan/factors-affecting-campus-placement>)

In [1]:

```
import pandas as pd
from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
```

Importing the dataset

In [2]:

```
data=pd.read_csv("Placement_Data_Full_Class.csv")
```

In [3]:

```
data.describe()
```

Out[3]:

	sl_no	ssc_p	hsc_p	degree_p	etest_p	mba_p	salary
count	215.000000	215.000000	215.000000	215.000000	215.000000	215.000000	148.000000
mean	108.000000	67.303395	66.333163	66.370186	72.100558	62.278186	288655.405405
std	62.209324	10.827205	10.897509	7.358743	13.275956	5.833385	93457.452420
min	1.000000	40.890000	37.000000	50.000000	50.000000	51.210000	200000.000000
25%	54.500000	60.600000	60.900000	61.000000	60.000000	57.945000	240000.000000
50%	108.000000	67.000000	65.000000	66.000000	71.000000	62.000000	265000.000000
75%	161.500000	75.700000	73.000000	72.000000	83.500000	66.255000	300000.000000
max	215.000000	89.400000	97.700000	91.000000	98.000000	77.890000	940000.000000

In [4]:

```
data.describe(include="all")
```

Out[4]:

	sl_no	gender	ssc_p	ssc_b	hsc_p	hsc_b	hsc_s	degree_p
count	215.000000	215	215.000000	215	215.000000	215	215	215.000000
unique	NaN	2	NaN	2	NaN	2	3	NaN
top	NaN	M	NaN	Central	NaN	Others	Commerce	NaN
freq	NaN	139	NaN	116	NaN	131	113	NaN
mean	108.000000	NaN	67.303395	NaN	66.333163	NaN	NaN	66.370186
std	62.209324	NaN	10.827205	NaN	10.897509	NaN	NaN	7.358743
min	1.000000	NaN	40.890000	NaN	37.000000	NaN	NaN	50.000000
25%	54.500000	NaN	60.600000	NaN	60.900000	NaN	NaN	61.000000
50%	108.000000	NaN	67.000000	NaN	65.000000	NaN	NaN	66.000000
75%	161.500000	NaN	75.700000	NaN	73.000000	NaN	NaN	72.000000
max	215.000000	NaN	89.400000	NaN	97.700000	NaN	NaN	91.000000



In [5]:

```
print(data.shape)
print(data.head())
```

(215, 15)

	sl_no	gender	ssc_p	ssc_b	hsc_p	hsc_b	hsc_s	degree_p	\
0	1	M	67.00	Others	91.00	Others	Commerce	58.00	
1	2	M	79.33	Central	78.33	Others	Science	77.48	
2	3	M	65.00	Central	68.00	Central	Arts	64.00	
3	4	M	56.00	Central	52.00	Central	Science	52.00	
4	5	M	85.80	Central	73.60	Central	Commerce	73.30	

	degree_t	workex	etest_p	specialisation	mba_p	status	salary
0	Sci&Tech	No	55.0	Mkt&HR	58.80	Placed	270000.0
1	Sci&Tech	Yes	86.5	Mkt&Fin	66.28	Placed	200000.0
2	Comm&Mgmt	No	75.0	Mkt&Fin	57.80	Placed	250000.0
3	Sci&Tech	No	66.0	Mkt&HR	59.43	Not Placed	NaN
4	Comm&Mgmt	No	96.8	Mkt&Fin	55.50	Placed	425000.0

Handling missing values

In [6]:

```
data.isnull()
```

Out[6]:

	sl_no	gender	ssc_p	ssc_b	hsc_p	hsc_b	hsc_s	degree_p	degree_t	workex	etest_p
0	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False
...
210	False	False	False	False	False	False	False	False	False	False	False
211	False	False	False	False	False	False	False	False	False	False	False
212	False	False	False	False	False	False	False	False	False	False	False
213	False	False	False	False	False	False	False	False	False	False	False
214	False	False	False	False	False	False	False	False	False	False	False

215 rows × 15 columns



In [7]:

```
data.isnull().sum()
```

Out[7]:

```
sl_no          0
gender         0
ssc_p          0
ssc_b          0
hsc_p          0
hsc_b          0
hsc_s          0
degree_p       0
degree_t       0
workex         0
etest_p        0
specialisation 0
mba_p          0
status         0
salary        67
dtype: int64
```

Replacing null salaries by 0

In [8]:

```
data['salary'].fillna(value=0 , inplace = True )
```

In [9]:

```
data.isnull()
```

Out[9]:

	sl_no	gender	ssc_p	ssc_b	hsc_p	hsc_b	hsc_s	degree_p	degree_t	workex	etest_p
0	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False
...
210	False	False	False	False	False	False	False	False	False	False	False
211	False	False	False	False	False	False	False	False	False	False	False
212	False	False	False	False	False	False	False	False	False	False	False
213	False	False	False	False	False	False	False	False	False	False	False
214	False	False	False	False	False	False	False	False	False	False	False

215 rows × 12 columns



In [10]:

```
data.isnull().sum()
```

Out[10]:

```
sl_no          0
gender         0
ssc_p          0
ssc_b          0
hsc_p          0
hsc_b          0
hsc_s          0
degree_p       0
degree_t       0
workex         0
etest_p        0
specialisation 0
mba_p          0
status         0
salary         0
dtype: int64
```

In [11]:

```
data.drop(['sl_no', 'ssc_b', 'hsc_b'], axis = 1 , inplace = True)
```

In [12]:

```
print(data.head())
```

	gender	ssc_p	hsc_p	hsc_s	degree_p	degree_t	workex	etest_p	\
0	M	67.00	91.00	Commerce	58.00	Sci&Tech	No	55.0	
1	M	79.33	78.33	Science	77.48	Sci&Tech	Yes	86.5	
2	M	65.00	68.00	Arts	64.00	Comm&Mgmt	No	75.0	
3	M	56.00	52.00	Science	52.00	Sci&Tech	No	66.0	
4	M	85.80	73.60	Commerce	73.30	Comm&Mgmt	No	96.8	

	specialisation	mba_p	status	salary
0	Mkt&HR	58.80	Placed	270000.0
1	Mkt&Fin	66.28	Placed	200000.0
2	Mkt&Fin	57.80	Placed	250000.0
3	Mkt&HR	59.43	Not Placed	0.0
4	Mkt&Fin	55.50	Placed	425000.0

Encoding

In [13]:

```
columns= ['gender','hsc_s','degree_t','workex','specialisation','status']
label_encoder = LabelEncoder()
for column in columns:
    data[column]= label_encoder.fit_transform(data[column])
data.head()
```

Out[13]:

	gender	ssc_p	hsc_p	hsc_s	degree_p	degree_t	workex	etest_p	specialisation	mba_p
0	1	67.00	91.00	1	58.00	2	0	55.0	1	58.80
1	1	79.33	78.33	2	77.48	2	1	86.5	0	66.28
2	1	65.00	68.00	0	64.00	0	0	75.0	0	57.80
3	1	56.00	52.00	2	52.00	2	0	66.0	1	59.43
4	1	85.80	73.60	1	73.30	0	0	96.8	0	55.50

Outlier detection

In [14]:

```
plt.figure(figsize = (15,10))

ax = plt.subplot(221)
plt.boxplot(data['ssc_p'])
ax.set_title('Secondary School Percentage')

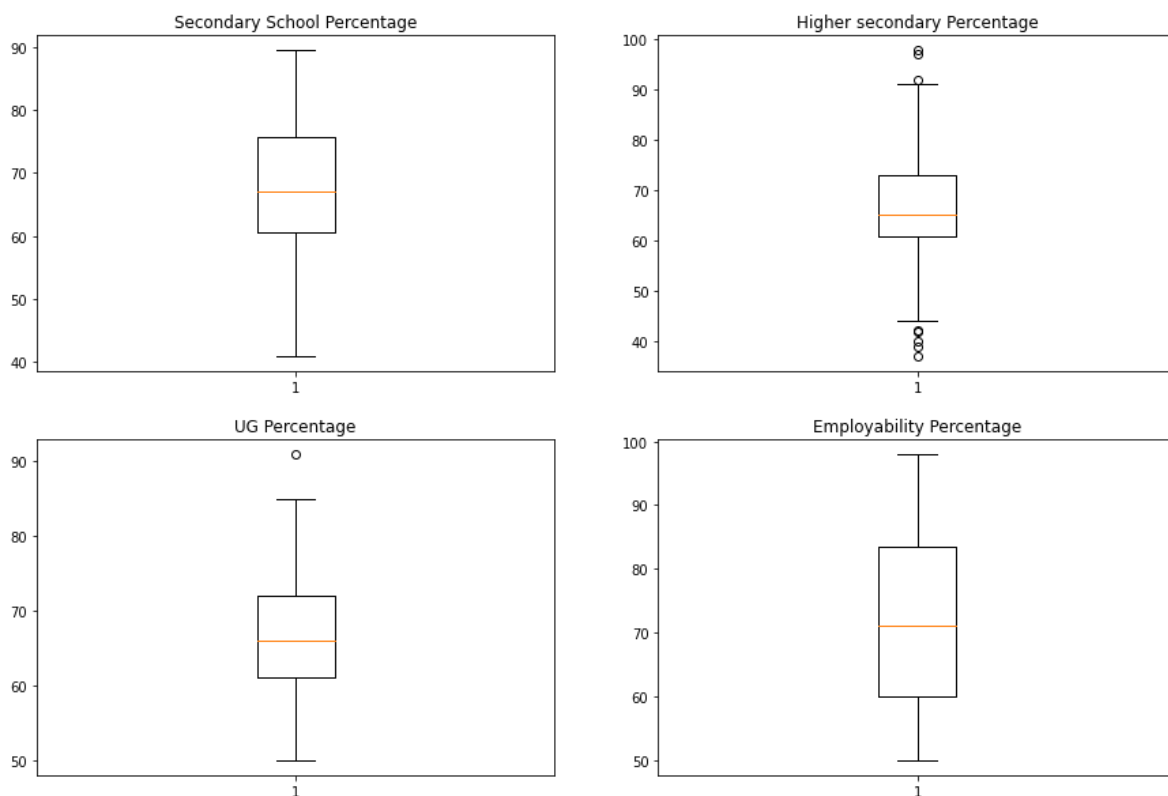
ax = plt.subplot(222)
plt.boxplot(data['hsc_p'])
ax.set_title('Higher secondary Percentage')

ax = plt.subplot(223)
plt.boxplot(data['degree_p'])
ax.set_title('UG Percentage')

ax = plt.subplot(224)
plt.boxplot(data['etest_p'])
ax.set_title('Employability Percentage')
```

Out[14]:

Text(0.5, 1.0, 'Employability Percentage')



In [15]:


```
Q1 = data['hsc_p'].quantile(0.25)
Q3 = data['hsc_p'].quantile(0.75)
IQR = Q3 - Q1

filter = (data['hsc_p'] >= Q1 - 1.5 * IQR) & (data['hsc_p'] <= Q3 + 1.5 * IQR)
data_filtered = data.loc[filter]

plt.boxplot(data_filtered['hsc_p'])
```

Out[15]:


```
{'whiskers': [<matplotlib.lines.Line2D at 0x1ba0236ddc0>,
<matplotlib.lines.Line2D at 0x1ba0237d160>],
'caps': [<matplotlib.lines.Line2D at 0x1ba0237d4c0>,
<matplotlib.lines.Line2D at 0x1ba0237d820>],
'boxes': [<matplotlib.lines.Line2D at 0x1ba0236dc40>],
'medians': [<matplotlib.lines.Line2D at 0x1ba0237db80>],
'fliers': [<matplotlib.lines.Line2D at 0x1ba0237dee0>],
'means': []}
```



Visualisation

In [16]:

```
plt.figure(figsize = (15,7))
plt.subplot(231)
ax = sns.countplot(x= 'gender' , data = data)
plt.subplot(232)
ax = sns.countplot(x= 'hsc_s' , data = data)
plt.subplot(233)
ax = sns.countplot(x= 'degree_t' , data = data)
plt.subplot(234)
ax = sns.countplot(x= 'specialisation' , data = data)
plt.subplot(235)
ax = sns.countplot(x= 'workex' , data = data)
plt.subplot(236)
ax = sns.countplot(x= 'status' , data = data)
```



Normalization

In [17]:

```
from sklearn.model_selection import train_test_split
x=data.iloc[:,0:10]
y=data.iloc[:,10]
x_train, x_test,y_train,y_test = train_test_split(x,y,test_size=.2,random_state=0)
# data normalization with sklearn
from sklearn.preprocessing import MinMaxScaler
import numpy as np
norm = MinMaxScaler().fit(x_train.iloc[:, np.r_[1,2,4,7,9]])
train_norm=norm.transform(x_train.iloc[:, np.r_[1,2,4,7,9]])
test_norm=norm.transform(x_test.iloc[:, np.r_[1,2,4,7,9]])
```


In [18]:

```
x_training = pd.DataFrame()
x_training['0']=x_train.iloc[:,0]
x_training['1']=pd.DataFrame(train_norm).iloc[:,0]
x_training['2']=pd.DataFrame(train_norm).iloc[:,1]
x_training['3']=x_train.iloc[:,3]
x_training['4']=pd.DataFrame(train_norm).iloc[:,2]
x_training['5']=x_train.iloc[:,5]
x_training['6']=x_train.iloc[:,6]
x_training['7']=pd.DataFrame(train_norm).iloc[:,3]
x_training['8']=x_train.iloc[:,8]
x_training['9']=pd.DataFrame(train_norm).iloc[:,4]
print(x_training)
```

	0	1	2	3	4	5	6	7	8	9
16	1	0.231620	0.620428	1	0.241463	0	1	0.128125	0	0.553973
135	0	0.956625	0.378913	2	0.571463	0	0	0.208333	1	0.377811
122	0	0.457818	0.411862	0	0.243902	0	1	0.270833	0	0.043853
22	0	0.219258	0.280066	2	0.268293	2	0	0.208333	1	0.353448
80	0	0.653004	0.356837	1	0.195122	0	1	0.125000	1	0.378186
..
67	1	0.609629	0.593081	1	0.365854	0	0	0.416667	0	0.228636
192	1	NaN	NaN	1	NaN	0	1	NaN	0	NaN
117	1	0.306007	0.494234	2	0.341463	2	0	0.208333	0	0.602324
47	1	0.240946	0.296540	1	0.153659	0	1	0.187500	0	0.507121
172	1	NaN	NaN	1	NaN	0	0	NaN	1	NaN

[172 rows x 10 columns]

In [19]:

```
print(x_training.isnull().sum()* 100 / len(x_training))
```

```
0      0.000000
1     20.930233
2     20.930233
3      0.000000
4     20.930233
5      0.000000
6      0.000000
7     20.930233
8      0.000000
9     20.930233
dtype: float64
```

In [20]:

```
# Even when the percentage of null values is less than 40%, we are not imputing the values
```

In [21]:

```
print(x_training.index[x_training.isnull().any(1)].tolist())
x_training.dropna(axis=0,inplace=True)
print(x_training)
```

```
[197, 173, 179, 187, 194, 210, 188, 200, 206, 208, 191, 182, 205, 212, 196,
213, 201, 183, 178, 199, 184, 214, 186, 176, 189, 180, 177, 175, 207, 202, 1
74, 204, 193, 195, 192, 172]
```

	0	1	2	3	4	5	6	7	8	9
16	1	0.231620	0.620428	1	0.241463	0	1	0.128125	0	0.553973
135	0	0.956625	0.378913	2	0.571463	0	0	0.208333	1	0.377811
122	0	0.457818	0.411862	0	0.243902	0	1	0.270833	0	0.043853
22	0	0.219258	0.280066	2	0.268293	2	0	0.208333	1	0.353448
80	0	0.653004	0.356837	1	0.195122	0	1	0.125000	1	0.378186
..
9	1	0.934071	0.263591	1	1.000000	0	0	0.194167	0	0.693403
103	1	0.333767	0.425535	2	0.238780	2	1	0.208333	1	0.227886
67	1	0.609629	0.593081	1	0.365854	0	0	0.416667	0	0.228636
117	1	0.306007	0.494234	2	0.341463	2	0	0.208333	0	0.602324
47	1	0.240946	0.296540	1	0.153659	0	1	0.187500	0	0.507121

[136 rows x 10 columns]

In [22]:

```
y_training=y_train.drop([197, 173, 179, 187, 194, 210, 188, 200, 206, 208, 191, 182, 205, 2
213, 201, 183, 178, 199, 184, 214, 186, 176, 189, 180, 177, 175, 207, 202, 174, 204, 193, 1
print(y_training)
```

```
16      1
135     1
122     1
22      1
80      1
..
9       0
103     1
67      1
117     1
47      1
```

Name: status, Length: 136, dtype: int32

In [23]:

```
x_testing = pd.DataFrame()
x_testing['0']=x_test.iloc[:,0]
x_testing['1']=pd.DataFrame(test_norm).iloc[:,0]
x_testing['2']=pd.DataFrame(test_norm).iloc[:,1]
x_testing['3']=x_test.iloc[:,3]
x_testing['4']=pd.DataFrame(test_norm).iloc[:,2]
x_testing['5']=x_test.iloc[:,5]
x_testing['6']=x_test.iloc[:,6]
x_testing['7']=pd.DataFrame(test_norm).iloc[:,3]
x_testing['8']=x_test.iloc[:,8]
x_testing['9']=pd.DataFrame(test_norm).iloc[:,4]
print(x_testing)
```

	0	1	2	3	4	5	6	7	8	9
198	0	NaN	NaN	1	NaN	1	0	NaN	1	NaN
37	0	0.262633	0.428336	2	0.243902	2	0	0.416667	1	0.074588
89	0	NaN	NaN	2	NaN	2	1	NaN	1	NaN
168	0	NaN	NaN	1	NaN	0	1	NaN	1	NaN
171	1	NaN	NaN	1	NaN	0	1	NaN	0	NaN
75	0	NaN	NaN	1	NaN	0	0	NaN	1	NaN
96	0	NaN	NaN	2	NaN	0	1	NaN	0	NaN
137	1	NaN	NaN	1	NaN	0	0	NaN	1	NaN
5	1	0.392756	0.411862	2	0.670732	2	1	0.500000	0	0.591829
83	1	NaN	NaN	2	NaN	2	1	NaN	0	NaN
55	1	NaN	NaN	2	NaN	0	0	NaN	1	NaN
145	1	NaN	NaN	2	NaN	2	0	NaN	1	NaN
160	1	NaN	NaN	2	NaN	2	1	NaN	1	NaN
112	1	NaN	NaN	1	NaN	0	0	NaN	1	NaN
74	1	NaN	NaN	1	NaN	0	0	NaN	0	NaN
203	1	NaN	NaN	1	NaN	0	0	NaN	1	NaN
126	0	NaN	NaN	2	NaN	2	1	NaN	0	NaN
12	0	1.000000	0.609555	2	0.365854	0	0	0.520833	1	0.790105
152	1	NaN	NaN	2	NaN	2	1	NaN	0	NaN

In [24]:

```
print(x_testing.isnull().sum()* 100 / len(x_testing))
```

```
0    0.00000
1    83.72093
2    83.72093
3    0.00000
4    83.72093
5    0.00000
6    0.00000
7    83.72093
8    0.00000
9    83.72093
dtype: float64
```

In [25]:

```
print(x_testing.index[x_testing.isnull().any(1)].tolist())
x_testing.dropna(axis=0,inplace=True)
print(x_testing)
```

```
[198, 89, 168, 171, 75, 96, 137, 83, 55, 145, 160, 112, 74, 203, 126, 153, 1
58, 169, 141, 209, 190, 144, 185, 86, 71, 63, 143, 97, 136, 162, 154, 90, 21
1, 106, 181, 139]
```

	0	1	2	3	4	5	6	7	8	9
37	0	0.262633	0.428336	2	0.243902	2	0	0.416667	1	0.074588
5	1	0.392756	0.411862	2	0.670732	2	1	0.500000	0	0.591829
12	0	1.000000	0.609555	2	0.365854	0	0	0.520833	1	0.790105
18	0	0.175884	0.362438	1	0.365854	0	0	0.750000	1	0.422414
15	0	0.320755	0.400824	1	0.167561	0	1	0.333333	0	0.265742
7	1	0.566255	0.428336	2	0.536585	2	1	0.125000	0	0.344828
33	0	0.642160	0.420099	2	0.268293	0	1	0.914792	0	0.667916

In [26]:

```
y_testing=y_test.drop([198, 89, 168, 171, 75, 96, 137, 83, 55, 145, 160, 112, 74, 203, 126,
print(y_testing)
```

```
37    1
5      0
12     0
18     0
15     1
7      1
33     1
```

Name: status, dtype: int32

Prediction using Logistic Regression

In [27]:

```
from sklearn.linear_model import LogisticRegression
LR = LogisticRegression()
LR.fit(x_training, y_training)
LR.score(x_testing, y_testing)
```

Out[27]:

0.7142857142857143

In [28]:

```
a=[[1,70,91,0,60,2,0,70,1,58.80]]
LR.predict(a)
```

C:\Users\SAKSHI NEERAJ\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
warnings.warn(

Out[28]:

array([1])

In [29]:

```
from sklearn.metrics import confusion_matrix
print(LR.predict(x_testing))
print(confusion_matrix(y_testing, LR.predict(x_testing)))
```

```
[1 1 1 0 1 1 1]
```

```
[[1 2]
```

```
 [0 4]]
```